Smart Sentiment Analytics: Decoding Social Media Emotions with Computational Intelligence

Nagamedha Sakhamuri
Department of Computer Science
Georgia State University
Atlanta, Georgia, United States
nsakhamuri1@student.gsu.edu

Nikitha Bonthala
Department of Computer Science
Georgia State University
Atlanta, Georgia, United States
bnagasatyadurganiki1@student.gsu.edu

Yanqing Zhang
Department of Computer Science
Georgia State University
Atlanta, Georgia, United States
yzhang@gsu.edu

***Abstract* - In today's digital era, people express emotions not only in conversations but also through tweets, comments, and online reviews. These unstructured, informal texts form an emotional footprint that holds valuable insights for businesses, platforms, and AI systems. This project presents a sentiment analysis engine tailored for real-world social media data, specifically Twitter and YouTube comments. We compare three modeling strategies: traditional TF-IDF with Logistic Regression, CNN-based classifiers, and DistilBERT transformers. Through extensive experimentation, we observed that DistilBERT significantly outperforms other models, achieving 87% accuracy and a 0.87 F1-score. Additional enhancements like emoji sentiment scoring and sarcasm detection were also explored to boost emotional understanding. The project lays a foundation for emotionally aware AI systems and opens pathways for multilingual and domain-specific sentiment applications.**

***Keywords* — Sentiment Analysis, BERT, Social Media Analytics, Emotion Detection, Deep Learning, NLP, Transformers, Machine Learning**

I. INTRODUCTION

In an increasingly digital world, users express emotions through tweets, reviews, and comments — often in informal, noisy formats filled with slang, emojis, and sarcasm. These emotional signals hold significant value for businesses, platforms, and AI systems aiming to understand public opinion and customer sentiment.

While traditional sentiment analysis models perform well on clean text, they struggle with the complexities of real-world social media data. This gap limits their effectiveness in applications like brand monitoring, content moderation, and user engagement.

To bridge this gap, our project presents a sentiment analysis system trained on real Twitter and YouTube data. We implemented and compared multiple models — from traditional TF-IDF with Logistic Regression to deep learning with CNN, and finally, a transformer-based DistilBERT model known for capturing contextual meaning. The goal was not only to improve classification accuracy but to build a robust, scalable foundation for emotionally aware AI systems that reflect how modern tools like ChatGPT or Grok interpret human emotion in text.

II. BACKGROUND & MOTIVATION

While sentiment analysis has become a common tool in natural language processing, most existing systems are trained on clean, well-structured text. In contrast, social media content is often short, noisy, and informal — filled with slang, abbreviations, emojis, and sarcasm.

This mismatch creates a gap between real-world expression and machine interpretation. Traditional models tend to misclassify or overlook subtle emotional cues, leading to inaccurate predictions in practical applications.

To address this, we focused on comparing different modeling approaches for classifying sentiment in real-world data. Our work explores whether newer transformer-based architectures like DistilBERT can outperform traditional techniques in handling the challenges of informal online text.

III. ARCHITECTURE

The core architecture of our sentiment analysis system is designed to mirror a real-world AI pipeline — moving from noisy user-generated input to refined, sentiment-classified output. The design is modular, interpretable, and easily extensible for future enhancements like sarcasm or emoji handling.

As shown in **Figure 1**, the architecture consists of five stages:

1. **Input Collection**: We gather real-world user content from Kaggle datasets — specifically, 1.6M tweets from **Sentiment140** and thousands of **YouTube comments**. These serve as diverse, unstructured sources to simulate actual internet discourse.

2. **Preprocessing**: Raw text undergoes cleaning steps including lowercasing, punctuation removal, emoji normalization, slang expansion, and

tokenization. This ensures consistency before feeding into downstream models.

3. **Core Engine Processing**: The heart of the architecture is a fine-tuned **DistilBERT** model. It processes the preprocessed text to generate contextual embeddings, which are then passed to a classification head.

4. **Output Generation**: The classification head assigns one of three sentiment labels — **Positive**, **Negative**, or **Neutral** — based on the learned contextual understanding.

5. **Future Add-ons**: The architecture is designed to support additional modules like **emoji sentiment scoring** or **sarcasm detection**, which can be integrated seamlessly in future iterations.

Figure 1: Sentiment Analysis Architecture

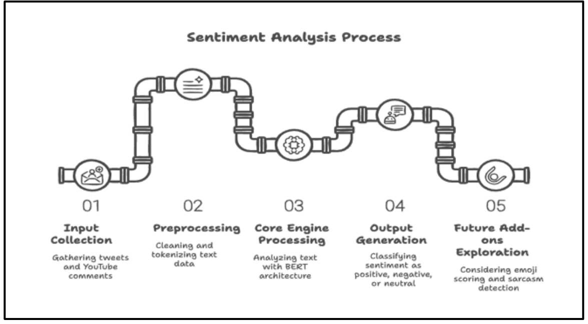
Sentiment Analysis Process
01
Input Collection
Gathering tweets and YouTube comments
02
Preprocessing
Cleaning and tokenizing text data
03
Core Engine Processing
Analyzing text with BERT architecture
04
Output Generation
Classifying sentiment as positive, negative, or neutral
05
Future Add-ons Exploration
Considering emoji scoring and sarcasm detection

This modular architecture allowed us to not only benchmark various modeling strategies — from traditional ML to transformers — but also maintain clarity and extensibility throughout development.

IV. METHODOLOGY

Our sentiment analysis system was developed through a multi-stage process involving data acquisition, preprocessing, and model development. Below we describe each component in detail.

Dataset Description: We utilized two diverse, real-world datasets from Kaggle to ensure model robustness across platforms and content styles:

- **Sentiment140_Dataset:** This dataset contains 1.6 million labeled tweets annotated using emoticons as weak sentiment indicators (positive, negative, neutral). The data is noisy and informal—perfect for testing sentiment models in real-world conditions.
    - https://www.kaggle.com/datasets/kazanova/sentiment140
- **US YouTube Comments Dataset**: This dataset includes user-generated YouTube comments, offering diverse vocabulary, use of slang, and frequent emojis. It was primarily used for cross-platform generalization.
    - https://www.kaggle.com/datasets/aashita/yt-comments-sentiment-dataset

These datasets together allowed us to train and evaluate models that could handle varied expressions, social media formats, and sentiment tones.

Preprocessing Pipeline: Raw text from social media is unstructured and messy. To prepare the data, we implemented the following preprocessing steps:

1. Lowercasing all text
2. Removing punctuation and stopwords
3. Emoji normalization (e.g., 😊 → positive)
4. Slang expansion (e.g., “idk” → “I don’t know”)
5. Tokenization using HuggingFace's tokenizer

These steps ensured cleaner and standardized input for both traditional and transformer-based models.

Figure 2: Preprocessing Steps

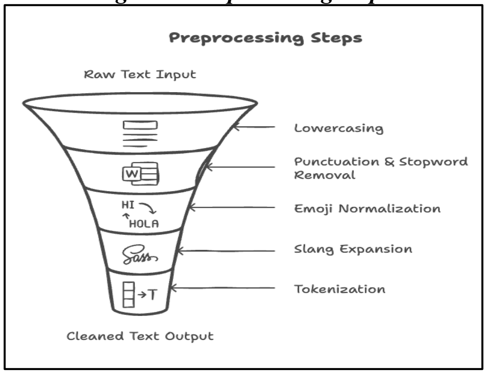

Modeling Techniques: We experimented with three types of models:

1. **Traditional Machine Learning Models:** We used TF-IDF features combined with models like Logistic Regression, Naive Bayes, and SVM. These models provided quick results and interpretability but struggled with slang, sarcasm, and informal tone.

2. **CNN-Based Deep Learning:** A 1D Convolutional Neural Network was implemented using word embeddings. It captured local patterns but failed to understand context and long-range dependencies in text.

3. **Transformer Model (DistilBERT):** We fine-tuned DistilBERT using HuggingFace Transformers. This model understood sentence-level context, handled emojis and slang better, and significantly outperformed others on all evaluation metrics. We trained DistilBERT in phases (1K, 10K, 50K samples), gradually improving performance.

   Figure 3 illustrates the DistilBERT architecture we leveraged — from tokenized input to embedding layers, transformer encoders, classification head, and sentiment output.

Figure 3: BERT Architecture

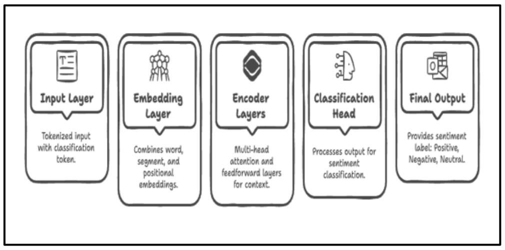
Input Layer
Tokenized input with classification token.
Embedding Layer
Combines word, segment, and positional embeddings.
Encoder Layers
Multi-head attention and feedforward layers for context.
Classification Head
Processes output for sentiment classification.
Final Output
Provides sentiment label: Positive, Negative, Neutral.

V. EVALUATION

To assess the performance of each model, we conducted experiments on multiple dataset sizes and evaluated them using standard classification metrics.

A. Evaluation Metrics

We used the following metrics for consistent comparison:

- **Accuracy**: Proportion of total correct predictions.
- **Precision**: Correctly predicted positives divided by total predicted positives.
- **Recall**: Correctly predicted positives divided by all actual positives.
- **F1 Score**: Harmonic mean of precision and recall.
- **Confusion Matrix**: Provides insight into class-wise prediction strengths and errors.

These metrics helped us go beyond accuracy and understand how each model performed on edge cases such as sarcasm, imbalanced sentiments, and emoji-based expressions.

B. Experimental Setup

To ensure fair benchmarking, we followed this phased evaluation plan:

TABLE 1: Experimental Setup

Phase	Dataset Size	Purpose
1	1,000	Initial benchmarking and debugging
2	10,000	Mid-scale testing of all models
3	50,000	Final DistilBERT fine-tuning and evaluation

All models were trained and tested on the same data splits using an 90:10 train-test ratio. DistilBERT training was performed on Google Colab with GPU acceleration using HuggingFace's Trainer framework. We also used additional visualization tools such as bar graphs, word clouds, and sentiment distribution plots to better interpret model behavior.

VI. RESULTS AND OBSERVATIONS

We evaluated three types of models across all dataset phases. The key findings and observations from each are summarized below.

A. Performance Comparison

As shown below, DistilBERT clearly outperformed both traditional and CNN models across all metrics. It showed better understanding of context, handled noisy inputs more effectively, and generalized well across both Twitter and YouTube data.

TABLE 2: RESULTS

Model	Accuracy	F1 Score
TF-IDF + Logistic Regression	72%	0.72
CNN	65%	0.65
DistilBERT	**87%**	**0.87**

B. Visualization Snapshots

- **Bar Charts**: Used to compare accuracy and F1-score across models.

Figure 4: Traditional CNN vs DistilBERT

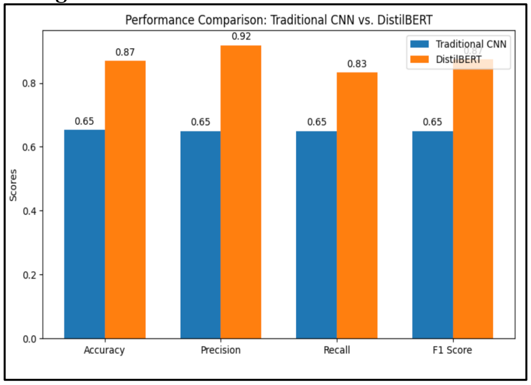

- **Confusion Matrices**: Provided insights into true vs. false predictions. Figure 5 shows the *DistilBERT confusion matrix image with best results — 87% accuracy.*

Figure 5: DistilBERT confusion matrix

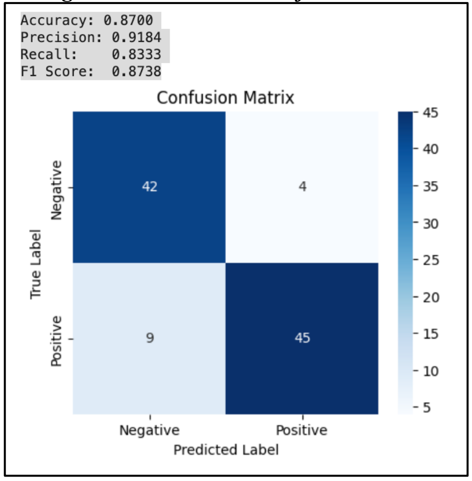

```
Accuracy: 0.8700
Precision: 0.9184
Recall:    0.8333
F1 Score:  0.8738
```

- **Sentiment Distribution**: Output prediction analysis showed that DistilBERT handled class balance effectively on both datasets.

Figure 6: **Sentiment Distribution**

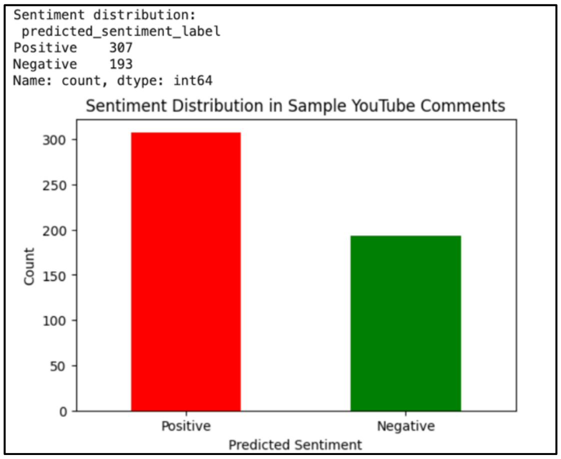

```
Sentiment distribution:
 predicted_sentiment_label
Positive    307
Negative    193
Name: count, dtype: int64
```

- **Word Clouds & TF-IDF Analysis:** Helped visualize most frequent and impactful terms for sentiment prediction.

Figure 7: Word cloud image of Positive & Negative Comments

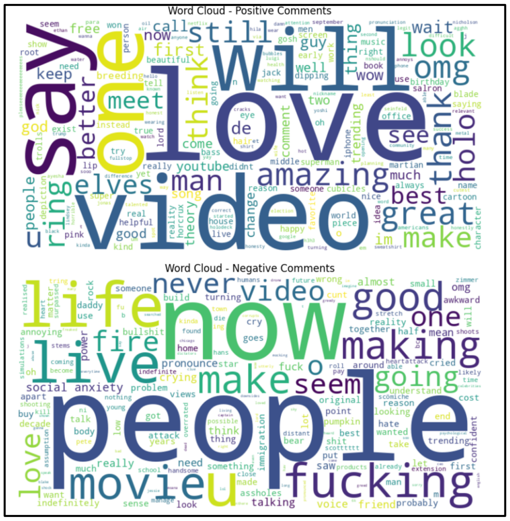

C. Key Observations

- Traditional models provided quick results but lacked depth in handling informal or sarcastic tone.
- CNN models performed slightly better with embeddings but failed to capture long-range dependencies.
- DistilBERT significantly improved accuracy and F1 score, particularly after training on a 50k sample and using data augmentation.

VII. CHALLENGES AND SOLUTIONS

Developing a sentiment analysis system for real-world social media content brought several challenges across data quality, model scalability, and evaluation. Below are the key challenges we faced and how we addressed them.

A. Noisy and Informal Text

➢ **Challenge:** Tweets and YouTube comments often include slang, emojis, hashtags, abbreviations, and sarcasm. These elements made it difficult for traditional models to interpret sentiment correctly.

➢ **Solution:** We implemented a robust preprocessing pipeline that included:

- Lowercasing
- Punctuation and stopword removal
- Emoji normalization
- Slang and abbreviation expansion (e.g., "idk" → "I don't know")
- Tokenization using HuggingFace's tokenizer

B. Transformer Model Resource Requirements

➢ **Challenge:** Training BERT-based models on larger datasets resulted in memory bottlenecks and longer training times.

➢ **Solution:** We used **DistilBERT**, a lightweight version of BERT, which required fewer resources. Additionally, we:

- Trained models in phases (1K → 10K → 50K samples)
- Utilized Google Colab with GPU support

C. Sarcasm and Emoji Interpretation

➢ **Challenge:** Sarcasm and emojis often distort literal meaning, making it difficult for models to capture the true sentiment.

➢ **Solution:** We explored:

- **Emoji sentiment scoring** through dictionary mapping
- A prototype **sarcasm detection** module (experimental phase)

These enhancements are functional and designed for future integration.

D. Evaluation Complexity

➢ **Challenge:** Accuracy alone was insufficient for evaluating model performance, especially with class imbalance and emotional subtleties.

➢ **Solution:** We used a comprehensive set of evaluation metrics: Accuracy, Precision, Recall, F1-Score, Confusion Matrix.

This provided deeper insight into model strengths and weaknesses.

E. Improving Generalization

➢ **Challenge:** With limited samples in early training phases, models risked overfitting and poor performance on unseen data.

- **Solution:** We applied **data augmentation** using techniques like **Random Word Swap**, increasing training diversity and improving robustness, especially in DistilBERT.

VIII. LIMITATIONS

LanguageScope: The current model only supports English. It may not generalize well to multilingual content or regional slang without further adaptation.

Real-Time Inference: The system operates in a batch-processing mode. It is not currently optimized for real-time comment analysis or live sentiment feedback.

Partial Integration of Add-ons: While we explored and prototyped emoji sentiment scoring and sarcasm detection, these modules are not yet fully integrated into the final model pipeline.

Domain-Specific Limitations: The model has not been fine-tuned for specific domains like healthcare, finance, or politics. It is trained on general public sentiment and may underperform on niche applications.

IX. FUTURE WORK

Multilingual Expansion: Extend sentiment classification capabilities to other languages such as Hindi, Telugu, and Tamil to support broader audience analysis.

Real-Time Monitoring System: Build an API-ready version for real-time sentiment analysis of tweets, product reviews, or live chat messages.

Add-on Enhancements: Fully integrate emoji sentiment scoring and sarcasm detection into the inference pipeline for improved emotion understanding.

Domain Fine-Tuning: Train on specific datasets in fields like customer support, political opinion mining, or healthcare to improve domain adaptability.

Inference Optimization: Explore faster BERT variants (e.g., TinyBERT, MobileBERT) and quantization techniques for faster and lighter deployment.

X. CONCLUSION

This project demonstrates the feasibility and effectiveness of building a sentiment analysis engine tailored for real-world social media content. By combining traditional machine learning, deep learning, and transformer-based models, we systematically evaluated performance across noisy, informal text.

Among the models tested, DistilBERT consistently outperformed others — achieving 87% accuracy and a 0.87 F1-score — while also handling slang, emojis, and context more effectively. Our phased experimentation, preprocessing pipeline, and modular architecture allowed us to adapt flexibly across datasets and evaluation stages.

In addition to strong results, this work lays the groundwork for future sentiment analysis systems that are multilingual, emotionally nuanced, and capable of real-time adaptation. We believe this project reflects the foundational logic behind modern AI tools like ChatGPT and Grok — systems that don't just understand words, but the emotions behind them.

REFERENCES

[1] J. Devlin, M. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," *arXiv preprint arXiv:1810.04805*, 2018.

[2] Hugging Face Transformers Library. [Online]. Available: https://huggingface.co/transformers/

[3] Kaggle Dataset – Sentiment140. [Online]. Available: https://www.kaggle.com/datasets/kazanova/sentiment140

[4] Kaggle Dataset – US YouTube Comments. [Online]. Available: https://www.kaggle.com/datasets/aashita/yt-comments-sentiment-dataset

[5] S. Ruder, "An overview of gradient descent optimization algorithms," *arXiv preprint arXiv:1609.04747*, 2016.

[6] A. Vaswani et al., "Attention is All You Need," *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[7] Scikit-learn: Machine Learning in Python. [Online]. Available: https://scikit-learn.org/

[8] NLTK – Natural Language Toolkit. [Online]. Available: https://www.nltk.org/

APPENDIX: RESOURCES

The complete implementation, models, and analysis notebooks are available at the following location:

Google Drive: Executed Code Notebooks, PPT, Software Manual:
https://drive.google.com/drive/folders/1uaNAeurbdx-1zIOPtFrP-ugq1_MPKTD6?usp=sharing